

LoRA adapter weights stay small on page one.

QLoRA quantization tradeoffs appear on page two.